

Phase 2 GT Final-OR Documentation

Gabriele Bortolato

March 27, 2023

1 Introduction

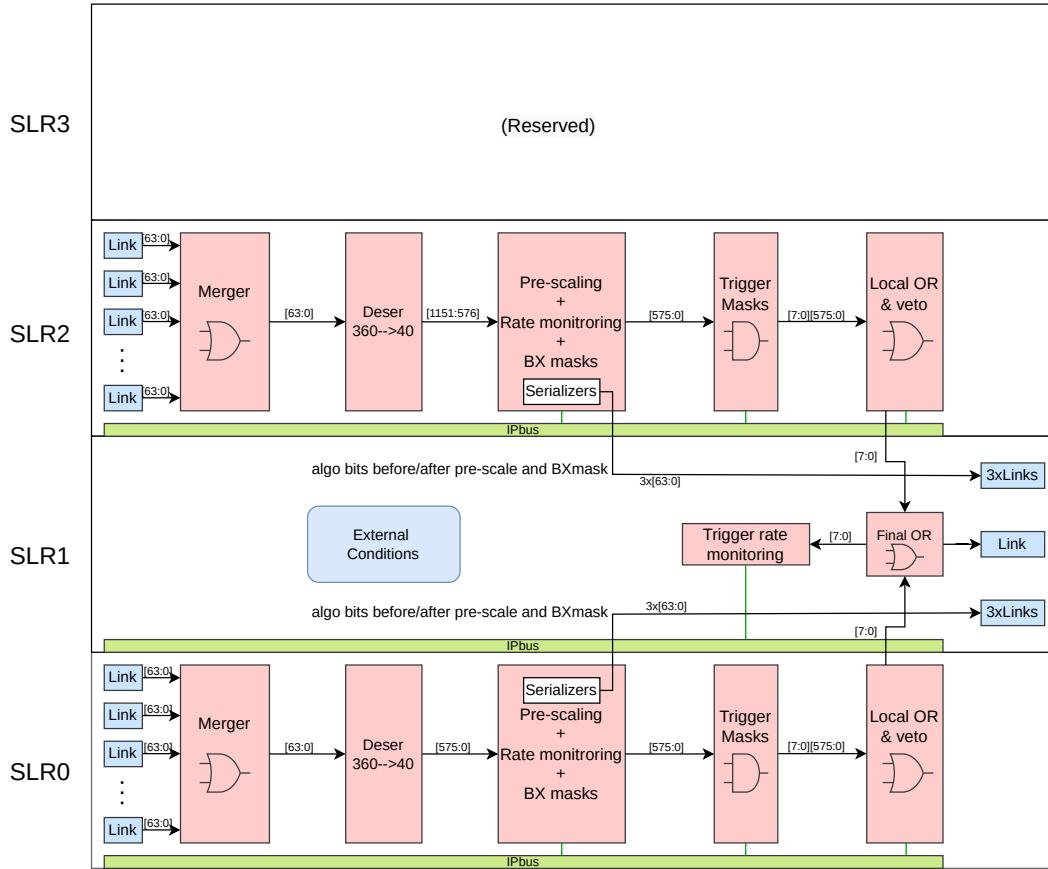


Figure 1: Caption

2 Modules

In this section main modules will be described in detail, starting from what the data see first up to the last OR. Only the modules within the EMP payload will be described, if more information about EMP framework is need visit the EMP website ([putref](#)) or the GitLab repository ([putref](#)).

2.1 Link merger

The link merger has to take as input the data links, coming from GT-Algorithm boards, and merge them into a single data link. To achieve this each metadata bits are OR-ed together, i.e. all the valid




Figure 2: Caption

bits are "merged" into a single bit. The data stream is merged in a similar way OR-ing together the data bits, but they are set to zero according to its valid bit this is done per link basis.

An issue that may arise is the fact that if one of the link is not aligned it will produce a wrong output. To tackle this issue a link mask is applied before the merging step, if a specific link is masked it will be set to a NULL value which means all metadata and data bits are set to 0. Such link mask value can be accessed via IPbus.

2.2 Deserializer




Figure 3: Caption

The deserializer module has to translate the 64-bit 360 MHz stream into 576-bit 40 MHz stream, to achieve this a frame counter starts count when link valid is asserted, it will count from 0 to 8 and start over.

At different frame counter values the 64-bit word is placed in different position within the 576-bit vector. A shift register is used to delay the incoming valid bit by exactly nine 360MHz clock cycles.

An alignment check is performed on the input metadata bits to verify if they behave as expected, e.g.

start-of-orbit signals happens when the frame counter is 0.

2.3 Pre-scaler

The pre-scale module developed by the Global Trigger group for the μ GT has been modified, but the main logic stayed the same [?]. The new module has been heavily optimized to reduce the resource foot print in order to reach the goal of 1152 total algorithms. Such optimization range from simple logic rearrangement to signal vectorization. The loading of the pre-scaler values is pictured in Fig. ?? where it is used a dual port ram to store the values, at each luminosity section the RAM is read and the values are stored in a register array, in the next luminosity section if the request update flag is asserted the factors are loaded inside the algo-slices.

Every algorithm has a pre-scaler with a pre-scale factor 24 bits wide¹. This module reduces the trigger-rate of the incoming algorithm by a given factor that can be set via software. For example a factor of 3 lets every third trigger pass blocking the first two.

A new feature introduced in 2019 is the possibility to set a fractional factor, each factor will be defined with an integer part and two digit fractional values. An example waveform with a factor of 1.75 is shown in Fig. 4.

The counter's increment of each algorithm assertion is 100, in fact with two fractional digit this corresponds to exactly 1. The factor of this example is 1.75, if the counter plus the increment is greater or equal to the pre-scale factor the algo_prescale flag is asserted and the counter is decreased by the pre-scale factor value.




Figure 4: Pre-scaler counter increment waveform.

2.4 Rate counter

The rate counters are used to extract the rate of each algorithm for monitoring purposes, in fact as was described in section ??, the level-1 trigger has a maximum accept rate and it is useful to have a monitor to spot potential rate mismatch with the expected values.

The counter is increased at each LHC clock cycle if the input algorithm is asserted while the reset condition is the start of a luminosity section or an external signal. Right before the reset its value is stored in a register and monitored by the experiment control room. In the implementation used in this work these rate counters are 32-bits wide.

Four rate counter are instantiated for each algorithm, one before the pre-scale, one after the pre-scale, one after the preview pre-scale and one post dead-time. The post dead-time counter is a slight variation of the other three, it increments the register only if the algorithm and the delayed L1A signal are both asserted, which means that only the algorithm assertions that can save the event are taking into account.

2.5 BX mask

The BX mask, as the name suggests, will exclude a given algorithm in a given BX number. The implementation is rather simple, in fact it is implemented as a dual port RAM, one port is connected to the programmable logic the other on is accessible via IPbus. This RAM has to be 3564 deep and

¹This number can be increased up to 32 bits

1152 wide, the first limitation is related to the number of BX in one orbit, the second limitation is related to the number of running algorithms. The IPbus suite can manage memory resource A32/D32 which means that $1152/32=36$ dualport ram in parallel².

The BX number is used as address and the content of this wide RAM is the mask of that specific BX number.

2.6 Algo slice

2.7 Local-OR and Trigger type masks

Unlike μ GT the Phase-2 GT will have multiple trigger bits which are selected to tackle different Physics, an example of 8 trigger bits is given in the table 1.

Index	Trigger type
0	Stadard Physics
1	Zero Bias
2	Minimum Bias
3	High Priority
4	Long-lived particle (LLP)
5	First event of a 2-event LLP
6	Second event of a 2-event LLP
7	Periodic validation event

Table 1: Example of 8 different trigger bits




Figure 5: Schematic view of the extraction of the i-esim local trigger bit.

To extract such trigger bits, first a mask is applied to the whole algorithm vector selecting only a fraction of the 576 bits³ if an algorithm act as veto it will be excluded from the final OR computation, then local OR⁴ is computed to each masks results, such outputs are the local trigger bits. The masks are loaded via an IPbus dual port RAM with the exact same structure of the pre-scale factors. A diagram of such process is shown in Fig. 5

2.8 Veto

Veto mask is simliar to the trigger masks the only difference is teh fact that it is need only one mask, the latter is loaded via IPbus to identify which algorithms will act as veto. If any of these algorithms is asserted the trigger with veto will be bound to zero, in hardware this is done with an AND gate.

2.9 Final-OR

The final OR is done combining the signals from the two monitoring SLRs. Each SLR outputs the local or (8-bits wide), the local or preview (8-bits wide) and the local veto.

2.10 Output interface

To be discussed

²18 per monitoring SLR

³This process is done in each monitoring SLR

⁴Local in the sens that is computed in each SLR




Figure 6: Veto logic.

3 Payload IPbus tree

4 Pre-scale column loading

Due to the high number of monitored algorithms, the pre-scale columns cannot be loaded from SW with generic IPbus register banks. In fact, multiple issues arise doing so, the obvious solution is to use dual port RAMs (DPRAM). The drawback of such approach is the time required to read the content of the DPRAM, in this case the column is split in two monitoring SLRs with 576 factors which means $\sim 14.4\mu s$ read time required.

In FW it are present two banks of registers, one to store the values from the RAM and one that is updated only at the beginning of a luminosity sections, a scheme of such structure is given in Fig. [ref](#)

Pre-scale column can be update only at the begin of a luminosity section, in this way a column is retained throughout the whole luminosity section. The pre-scale column loading follow these steps:

- A new column is loaded into the IPbus RAM via SW
- A load flag is pulsed via SW
- If a load flag is sensed in the HW a FSM [put ref](#) it will trigger the start of the loading routine

The loading routine is depicted in the Fig. [put ref](#). The Check state evaluates if the column can be loaded safely in the register bank, avoiding any possible partial loading, in fact if a read routine starts right before the begin of a luminosity section the factors towards the end are not updated. The check is done on the orbit number, if this number is not the last orbit or the one before it the reading procedure can take place otherwise it will wait the next luminosity section.

4.1 SW code example

Random pre-scale factors set in SW example:

```
# Pre-scale default factors
prsc_fct      = np.uint32(100 * np.ones((2,576))) # 1.00
```

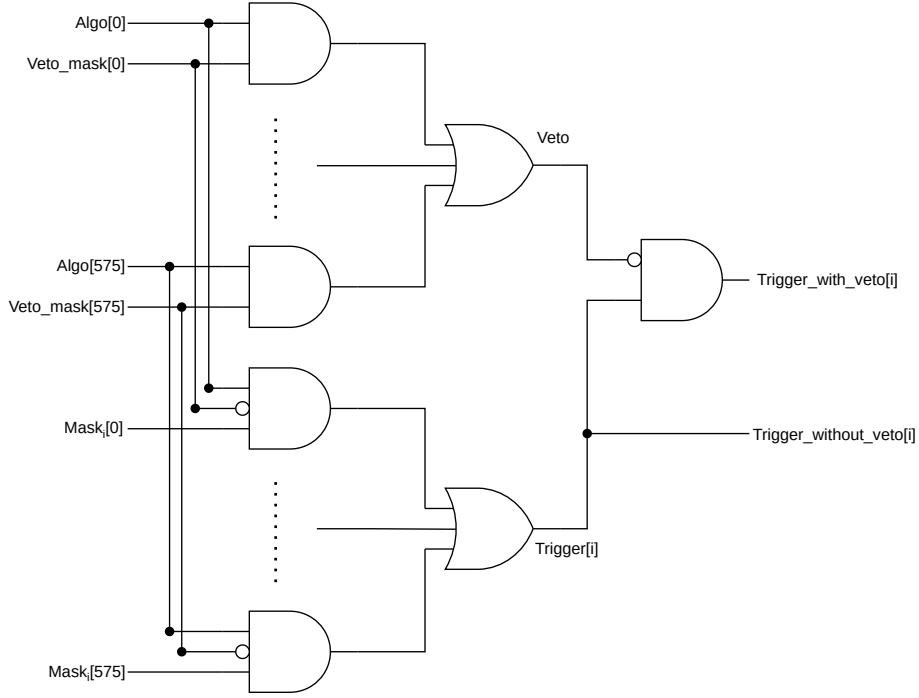


Figure 7: Trigger and veto mask logic.

State	Change conditions	Output
Idle	Load flag = '1'	Request_update='0' addr = 0
Check	$\text{orbit_nr}[\text{LS_BIT}-1:0] \leq 2^{\text{LSBIT}} - 2$	Request_update='0' addr = 0
Load	$\text{addr} \geq \text{RAM_DEPTH}-1$	Request_update='0' addr = addr+1
Done	Only one clock cycle	Request_update='1' addr = 0 LS_mark = LS_nr+1

Table 2: Caption

```

prsc_fct_prvw = np.uint32(100 * np.ones((2,576))) # 1.00
# Spit indexes in two SLR subsets
index_low = index[np.where(index < 576)[0]]
index_high = index[np.where(index >= 576)[0]]
#random assign
prsc_fct[1][np.int16(index_high - 576)] =
    np.uint32(np.random.randint(100, 2 ** 24, len(index_high)))
prsc_fct[0][np.int16(index_low)] =
    np.uint32(np.random.randint(100, 2 ** 24, len(index_low)))
#load the data into the IPbus RAM
HWtest.load_prsc_in_RAM(prsc_fct, PRESCALE_INDEX)
#same routine for the preview
prsc_fct_prvw[1][np.int16(index_high - 576)] =
    np.uint32(np.random.randint(100, 2 ** 24, len(index_high)))
prsc_fct_prvw[0][np.int16(index_low)] =

```

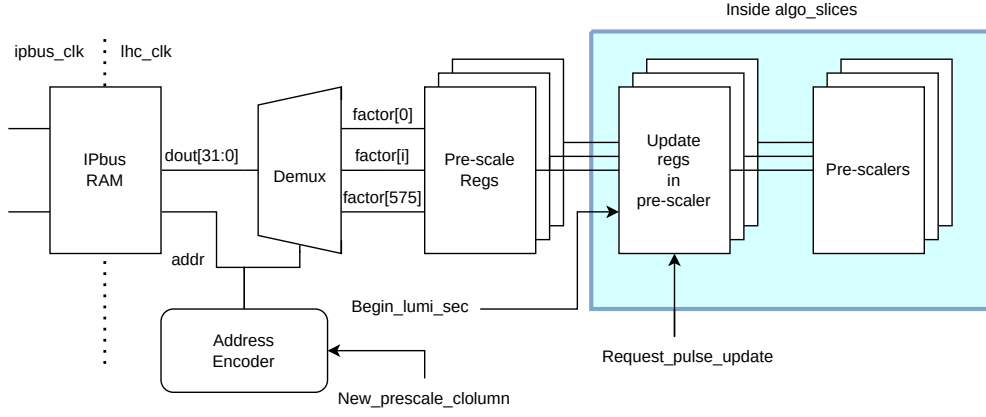


Figure 8: Caption

Reg name	32-bit Regs	RAM Depth	RAM Width
Pre-scale	576	1024	32
Pre-scale preview	576	1024	32
Trigger type masks	$\frac{576}{32} \times 8 = 144$	256	32
Veto mask	$\frac{576}{32} \times 1 = 18$	32	32
BX masks	$\frac{576}{32} \times 3564 = 65772$	4096	576
Algo rate counters ($\times 4$)	576	1024	32
Trigger rate counters ($\times 8$)	8	8	32

Table 3: Caption

```
np.uint32(np.random.randint(100, 2 ** 24, len(index_low)))
HWtest.load_prsc_in_RAM(prsc_fct, PRESCALE_PREVIEW_INDEX)
```

New column flag SW example:

```
#Send the new column flag
HWtest.send_new_prescale_column_flag()
time.sleep(2)
#Read the Lumi section mark
ls_prescale_mark = HWtest.read_lumi_sec_prescale_mark()
print("SLR 2 pre-scale column loaded mark = %d" % ls_prescale_mark[0])
print("SLR 3 pre-scale column loaded mark = %d" % ls_prescale_mark[1])
```

5 Trigger types masks and veto loading

The trigger type masks and the veto mask work in a similar fashion, but the RAM depth changes in relation to the bits that needs to be sent, a summary of registers required per SLR is given in the table [put ref.](#)

6 Rate counters values dumping

7 Final Or standalone test

A set of tests have been developed to check the Final-OR design correct behaviour. Such tests are unitary test in the sense that on a specific feature is tested each time. The tests developed are the following:

- **Pre-scaler test:** It will check if the pre-scalers and pre-scalers preview works as expected

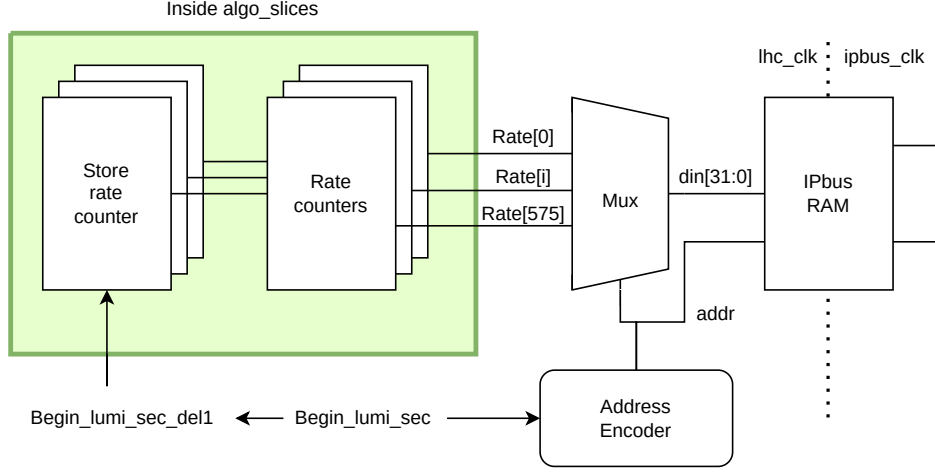


Figure 9: Caption

- **Trigger type masks test:** Here the trigger type masks logic is evaluated.
- **Veto mask test:** Similar to the Trigger type mask, but in this case only the veto mask is tested.
- **BX mask test:** test to evaluate the correct behaviour of the BX mask.

In the following sections each test is explained in detail.

7.1 Pre-scaler test

The EMP framework allows the injection of patternfiles directly on the datapath buffers, such buffers are limited in depth, in fact it is possible to inject up to 1024 frames, or 113^5 packets with TMUX=1.

As was mentioned in section [put ref](#) the Final-OR can manage up to 1152 algorithms. To produce a patternfiles it is necessary to first decide where (in terms of BXs) and which algorithm will fire, to do so a matrix 1152x113 is produced. The entries of such matrix can be either $[True, 1]$ or $[False, 0]$, the user can select the probability that an algorithm fires in a specific BX, this probability is constant throughout the whole matrix.

$$Matrix_{algo} = \underbrace{\begin{bmatrix} 0 & \dots & 1 & \dots & 1 \\ \vdots & \ddots & \dots & 0 & \vdots \\ 0 & \dots & 0 & \dots & 1 \\ \vdots & 1 & \dots & \ddots & \vdots \\ 0 & \dots & 0 & \dots & 1 \end{bmatrix}}_{113} \left. \vphantom{\begin{bmatrix} 0 & \dots & 1 & \dots & 1 \\ \vdots & \ddots & \dots & 0 & \vdots \\ 0 & \dots & 0 & \dots & 1 \\ \vdots & 1 & \dots & \ddots & \vdots \\ 0 & \dots & 0 & \dots & 1 \end{bmatrix}} \right\} 1152$$

Once the matrix is produced a patternfiles is generated starting from the position of such algorithm indexes distributing the isgnals in the links pool randomly.

The EMP framework will repeat the pattern files once every orbit, wich means that durement a **luminosity section** the pattern played 2^{LSBIT} times. The next step is to extract the expected unprscaled counts for each algorithm, to do so the columns are summed together resulting in a 113-wide vector.

⁵ $int(\frac{1024}{9}) = 113$, the link clock is 360MHz, in a BX 9 clock cycles are present

$$\begin{bmatrix} 0 & \cdots & 1 & \cdots & 1 \\ \vdots & \ddots & \cdots & 0 & \vdots \\ 0 & \cdots & 0 & \cdots & 1 \\ \vdots & 1 & \cdots & \ddots & \vdots \\ 0 & \cdots & 0 & \cdots & 1 \end{bmatrix} \rightarrow \vec{R} = \begin{bmatrix} 10 \\ \vdots \\ 32 \\ \vdots \\ 22 \end{bmatrix} \left. \vphantom{\begin{bmatrix} 0 & \cdots & 1 & \cdots & 1 \\ \vdots & \ddots & \cdots & 0 & \vdots \\ 0 & \cdots & 0 & \cdots & 1 \\ \vdots & 1 & \cdots & \ddots & \vdots \\ 0 & \cdots & 0 & \cdots & 1 \end{bmatrix}} \right\} 1152$$

The counts before pre-scale are computed with the formula:

$$Counts = \vec{R} \cdot 2^{LSBIT} \quad (1)$$

Then a pre-scale column is loaded into the firmware, such column can be either randomly generated or with equally spaced values starting from 1.00 to $\frac{2^{24}-100}{100} = 167771.16$. The theoretical counts after the pre-scalers are evaluated with the formula:

$$Counts = (\vec{R} \cdot 2^{LSBIT}) \oslash \vec{C} \quad (2)$$

Finally the firmware is loaded into the board and the rate counters are read and compared against the theoretical ones, if the mismatch is greater than 1 an error is raised and the test is failed.

7.2 Trigger type test

The trigger type masks test works in a similar way, but the final or counts are checked. To do so the algo matrix is split into 8 sub matrices maintaining the 113 width, which means that 8 subsets of algorithms are selected from the pool.

From these selections 8 different trigger type masks are extracted and saved to file to be passed to the RateChecker script.

Finally the column-wise OR is computed to such matrices resulting in a 8 113-wide vectors, the last step is to compute the element sum of such vectors resulting in 8 different counts for each trigger type. Then the results are compared against the hardware extracted ones.

put image

7.3 Veto mask test

Starting from a generated algo matrix a number of veto algorithms are extracted and a veto mask is produced, then the Final-OR counts are computed with and without the veto. In this test all trigger type masks are set as pass through⁶.

7.4 BX mask test

The test starts again with an algorithm matrix, then a BX mask matrix is produced in a similar way, in this case the probability of assertion is much greater (40%). Then the two matrices are bit-wise AND together to get the algorithm BXmasked matrix. From the algorithm matrix is extracted the input patternfile to sent to the board, from the BXmask matrix is extracted the relative BX masks and from the algorithm BXmasked matrix counts before pre-scalers are evaluated. Such counts are then checked against the HW results.

7.5 Suppress trigger during calibration test

Not yet defined

⁶As it was mentioned these are unitary tests, so they will test only one feature